

Machine Learning Models Comparison

Model	Components and Purposes	Primary Loss Function	Attention Mechanism	Pre-training Objective	Input Transformation Method
BERT (text)	<ul style="list-style-type: none"> Transformer Encoder: Contextual encoding 	Cross-entropy	Bidirectional self-attention	Masked language modeling and Next sentence prediction	Learned token embeddings + positional encodings
GPT (text)	<ul style="list-style-type: none"> Transformer Decoder: Autoregressive generation 	Cross-entropy	Causal self-attention	Next token prediction	Learned token embeddings + positional encodings
BART (text)	<ul style="list-style-type: none"> Transformer Encoder: Encoding corrupted input Transformer Decoder: Autoregressive reconstruction 	Cross-entropy	Bidirectional (encoder) + Causal (decoder) self-attention	Denoising autoencoding tasks (e.g., text infilling, sentence permutation)	Learned token embeddings + positional encodings
T5 (text)	<ul style="list-style-type: none"> Transformer Encoder: Task-specific encoding Transformer Decoder: Task-specific generation 	Cross-entropy	Encoder-decoder attention + Self-attention	Span corruption	Learned token embeddings + relative positional encodings
CLIP (text/vision)	<ul style="list-style-type: none"> Vision Transformer/ResNet: Image encoding Text Transformer: Text encoding 	Contrastive loss	Self-attention (in Transformers)	Image-text alignment	Image: Patch embeddings or CNN Text: Token embeddings
DALL-E (text/vision)	<ul style="list-style-type: none"> Discrete VAE: Image tokenization Transformer: Text-to-image generation 	Cross-entropy	Causal self-attention	Autoregressive image generation conditioned on text	Text: Token embeddings Image: dVAE to discrete tokens

Model	Components and Purposes	Primary Loss Function	Attention Mechanism	Pre-training Objective	Input Transformation Method
Diffusion Models (vision)	<ul style="list-style-type: none"> • U-Net (often): Feature extraction and denoising • Noise predictor: Estimates noise to be removed • Conditioner: Guides generation process (optional) 	Mean Squared Error (typically)	Self-attention (in U-Net) Cross-attention (for conditioning, if used)	Denoising score matching	Direct operation on pixel space + noise level embedding + optional conditioning signal
GAN (vision)	<ul style="list-style-type: none"> • Generator: Data generation • Discriminator: Real/fake classification 	Adversarial (minimax)	None (in basic GANs)	Distribution matching	Generator: Dense/Convolutional layers from noise vector Discriminator: Convolutional or dense layers
VQ-VAE (vision/audio)	<ul style="list-style-type: none"> • Encoder: Input encoding • Vector Quantizer: Discrete representation • Decoder: Reconstruction 	Reconstruction loss + Vector Quantization loss	None (typically)	Discrete representation learning	Convolutional encoder to latent space, then vector quantization
SimCLR (vision)	<ul style="list-style-type: none"> • Data augmentation: Create positive pairs • Base encoder (e.g., ResNet): Feature extraction • Projection head: Representation for contrastive loss 	Contrastive loss (e.g., NT-Xent)	None	Self-supervised visual representation learning	Convolutional layers (ResNet) + MLP projection head

Model	Components and Purposes	Primary Loss Function	Attention Mechanism	Pre-training Objective	Input Transformation Method
Wav2Vec (audio)	<ul style="list-style-type: none"> • Feature encoder: Raw audio embedding • Context network: Contextual representation • Quantizer: Discrete unit learning (in Wav2Vec 2.0) 	Contrastive loss	Self-attention (in Wav2Vec 2.0)	Self-supervised audio representation learning	Convolutional layers for raw audio embedding
Whisper (audio)	<ul style="list-style-type: none"> • Encoder: Audio feature extraction • Decoder: Text generation 	Cross-entropy	Encoder-decoder attention + Self-attention	Speech recognition and translation	Convolutional layers for audio, learned token embeddings for text